

MERIDIAN JOINT COMMAND · DIGITAL ASSURANCE

MERIDIAN VANTAGE


Hybrid Deployment Orchestrator
Cloud + On-Prem + Air-Gapped simulation

One model. Three perimeters. Every decision justified. A simulator and control room that deploys the same signed AI model everywhere, routes each job by classification policy, refuses what must not run, and proves it in a hash-chained ledger.

TEAM	Event Horizon
REPOSITORY	github.com/copsman/Hybrid-Deployment-Orchestrator
LIVE DEMO	https://hybrid-deployment-orchestrator.vercel.app
RUN LOCALLY	<code>npm ci && npm run build && npm start → http://localhost:3000 (production) · npm run demo · npm test</code>

NEXT.JS 16 · TYPESCRIPT · REACT-THREE-FIBER · SHADCN/UI · ED25519 · SHA-256

VITEST · FAST-CHECK · PLAYWRIGHT · VERCEL



CONTROL ROOM AFTER THE SCENARIO · 3 JOBS ROUTED · 3 REFUSED ·
PARITY VERIFIED · LEDGER INTACT

CLOUD
Meridian Cloud
Region North · elastic
· external

ON-PREM
Fort Meridian
Datacentre · 4 GPU ·
proxied

AIR-GAPPED
Enclave OBSIDIAN · 2
GPU · no network ·
diode

Government AI never lives in one place. The model still has to.

THE PROBLEM

Some data may go to public cloud, some must stay on-prem, and the most sensitive must never leave an air-gapped enclave. Today that means three pipelines maintained by hand, a change board deciding by email where a request may run, spreadsheets tracking which model version sits where, and artefacts carried into the enclave on media nobody verifies. Each step is a place where a SECRET request lands in the wrong perimeter or the enclave runs a stale or tampered model.

WHO BUYS AND DEPLOYS IT · WHAT IT REPLACES

- **Buyer:** the digital-assurance or platform directorate of a defence or government organisation that already runs a sovereign cloud tenancy, a datacentre and accredited enclaves, and the integrators who build for them.
- **Operator:** the platform team that owns all three environments and the security officers who approve what enters the enclave.
- **Replaces:** per-environment pipelines, manual routing decisions, unverified media transfers and audit spreadsheets, with one policy plane over the platforms they already run (Kubernetes, vLLM, registries, a diode).

WHY THIS APPROACH NOW, NOT THE ONE OF THREE YEARS AGO

- **Open-weight models** served by vLLM behind an OpenAI-compatible gateway make "the same model in every perimeter" practical; three years ago the capable models were API-only.
- **Policy as code** (OPA, Cedar) and **supply-chain signing** (Sigstore, SLSA) are standard practice; the router and the signed pipeline apply the same ideas to model deployment.
- **Sovereign and air-gapped cloud** and unidirectional gateways are commercial products; regulation increasingly demands a tamper-evident record of automated decisions.

OBJECTIVE · THE MINIMUM DELIVERABLE, AND HOW WE MEET IT

BRIEF ASKS FOR	MERIDIAN // VANTAGE DELIVERS
Three simulated environments with different properties	Cloud (elastic, metered, external), On-prem (4 fixed GPU slots, queue, proxied egress), Air-gapped (2 slots, no network, diode import)
Classification scheme and routing policies	OPEN / RESTRICTED / SECRET / ONYX plus attributes (personal data, live retrieval, releasability, model version); eight rules as data, default DENY
Router that logs decision and justification	Every submission yields a full record: rule trace, candidate table, verdict, tie-break, ledger hash
Deployment view incl. how the artefact reached the enclave	Version matrix with SHA-256 parity; pipeline view from build to diode, quarantine, two-person approval, verify, load
Blocked-request path with the reason shown	Three distinct refusals: no accredited environment, version not yet imported, contradictory constraints; plus what-if counterfactuals

DESIGN PRINCIPLES

- **Deny by default.** A job runs only where a rule explicitly grants it; capability checks can only remove options, never add them.
- **Nothing spills downward.** A full on-prem queues; it never bursts RESTRICTED work to cloud.
- **Refuse rather than downgrade.** If the enclave lacks the requested version, the job is refused with the reason, not silently run on an older model.
- **Everything is evidence.** Decisions and imports are links in a hash chain that anyone can recompute.
- **Proves itself headless.** The engine is pure TypeScript; the film on top is optional.

One policy plane, three real stacks, one signed artefact through the diode.

JOB

title, classification, personal data, live retrieval, releasability, model version, latency · validated by zod, never crashes



POLICY ENGINE

rules.json evaluated top-down: RESTRICT_TO grants, EXCLUDE subtracts, REFUSE ends · default DENY · every rule traced



CAPABILITY + SELECTION

CAP-VERSION (loaded there?), CAP-EGRESS (network?), CAP-SLOTS (queue) · then free capacity → cost → free slots



ENVIRONMENT

dispatch or queue in cloud / on-prem / enclave; enclave egress guard blocks any outbound call



LEDGER


$\text{hash}_n = \text{SHA-256}(\text{hash}_{n-1} \parallel \text{canonical-json}(\text{entry}_n))$ · decisions, dispatches, imports · verifiable, tamper-evident

ARTIFACT PIPELINE BUILT → SIGNED (Ed25519 over canonical manifest) → PUBLISHED (cloud registry) → MIRRORED (on-prem Harbor via proxy) → STAGED → **IN-DIODE** (chunked, one-way, no ACK; return attempts bounce and are counted) → QUARANTINE (AV + file type) → **VERIFYING** (two distinct operators) → IMPORTED (signer = pinned key · signature valid · SHA-256 matches · size · not a downgrade) → LOADED **REJECTED on any failed check**

SAME MODEL · THREE STACKS

LAYER	CLOUD	ON-PREM	AIR-GAPPED
App	Vercel (Next.js)	Next.js standalone on k3s	Same image, diode-imported
AI gateway	Vercel AI Gateway	Self-hosted OpenAI-compatible gateway	Same, offline, no external providers
Model serving	Managed GPU inference (vLLM)	vLLM on 4 GPU nodes	vLLM offline, weights from signed bundle
Database	Supabase Cloud (Postgres, pgvector)	Self-hosted Supabase	Postgres + pgvector, local backups
Identity	Supabase Auth	Keycloak	Keycloak offline, smart-card
Registry	GHCR	Harbor mirror (proxy sync)	Harbor, fed only by verified imports
Secrets · Observability	Vercel env · Langfuse Cloud	Vault · Langfuse + Prometheus	Vault sealed · Langfuse offline
Updates · Network	CI/CD push · external	Scheduled proxy pull · proxied	Diode import, two-person · none

Parity compares the SHA-256 of the loaded weights across the three environments. In production: evaluator → OPA/Cedar sidecar, signer → Sigstore cosign, diode → commercial unidirectional gateway.



DIRECTOR MODE: CHUNKS CROSS THE DIODE; THE ENCLAVE'S ACK BOUNCES OFF THE GATE

UNDER THE HOOD

- Pure TypeScript engine** runs in the browser, in Node (npm run demo) and in the tests; deterministic seed, clock and PRNG.
- Crypto:** SHA-256 + Ed25519 (RFC 8032, strict) from audited noble libraries over canonical JSON. No encryption, so no nonce to reuse.
- Control room:** Next.js 16, shadcn/ui, a react-three-fiber scene with a 2D fallback, and a scripted director that plays the scenario like a film.
- What-if:** dry-run plus single-change counterfactuals for refused jobs.

Three classifications, three environments, three distinct refusals.

THE DEMO SCENARIO · IDENTICAL ON EVERY RUN (SEED)


JOB	CLASSIFICATION · ATTRIBUTES	OUTCOME	WHY
A	OPEN public-affairs brief	→ CLOUD	POL-05 permits cloud or on-prem; cloud wins on cost (6.3 vs 10.2 credits)
B	RESTRICTED personal data	→ ON-PREM	POL-04 sovereign perimeter; POL-06 excludes cloud for personal data
C	SECRET needs model 1.4.0	REFUSED	CAP-VERSION: enclave still runs 1.3.0; refuse rather than downgrade
—	Ship 1.4.0 through the pipeline	PARITY OK	Same SHA-256 in cloud, on-prem and enclave after diode + two-person import
C'	SECRET same job again	→ AIR-GAPPED	POL-02: SECRET executes only inside the enclave
D	ONYX compartmented	REFUSED	POL-01: no environment in this deployment is accredited for ONYX
E	SECRET live external retrieval	REFUSED	POL-03: SECRET means the enclave; the enclave has no network

AUTOMATED CHECKS

- **40 vitest tests**, property-based with fast-check: SECRET never on cloud or on-prem; ONYX always refused; live retrieval never in the enclave; personal data never in cloud; every verdict has a rule id and justification; ledger tampering detected at any position; flipped byte or wrong signer always rejected; schema never throws.
- **Headless demo** exits non-zero on failure; CI runs it twice and diffs decisions, digests and ledger head.
- **Playwright** plays the UI to "SCENARIO COMPLETE", breaks and restores the ledger, submits garbage JSON, resets.
- **Fresh clone**: `npm ci && npm run build && npm start` serves the production build; demo, tests and build pass from a clean checkout.

UNEXPECTED INPUT AND SECOND RUNS

- Malformed JSON, unknown classifications, wrong types and extra fields are rejected with field-level messages; nothing is dispatched.
- Downgrade imports are refused; tampered bundles are rejected with the failing check named.
- Reset restores the seed: a second run reproduces the first, hash for hash.
- No WebGL or a scene error drops to the 2D map; an error boundary keeps the screen readable.



DECISION RECORD: VERDICT, JUSTIFICATION, RULE TRACE, CANDIDATES, LEDGER HASH



IMPORT CEREMONY: SCAN, TWO APPROVALS, PINNED KEY, SIGNATURE, DIGEST, LOAD

CRYPTOGRAPHY, APPLIED DELIBERATELY

SHA-256 for digests and the ledger chain; Ed25519 (strict RFC 8032) over a canonical JSON encoding that is property-tested for key-order independence, so a re-serialised manifest verifies and an edited one does not. Nothing is encrypted: integrity and authenticity are the requirement, confidentiality across the gap is the diode's physical property, so there is no nonce to reuse. Keys derive from the demo seed; in production the build key lives in an HSM and only the public half is pinned in each environment.

The prototype completes the brief end to end, deterministically, from a single command.

3 / 3JOBS ROUTED TO THE
CORRECT PERIMETER**3**REFUSALS, EACH WITH
A DISTINCT LOGGED
REASON**2**MODEL VERSIONS WITH
DIGEST PARITY IN ALL
THREE ENVS**33**HASH-CHAINED LEDGER
ENTRIES PER RUN,
VERIFIABLE**40 + 2**UNIT / PROPERTY
TESTS + BROWSER
TESTS, ALL GREEN**0**API KEYS, NETWORK
CALLS OR DATABASES
REQUIRED

WHAT IS NEW HERE

- **Refuse, then explain.** What-if counterfactuals tell an operator the one change that would make a refused job routable, without dispatching anything.
- **Decisions as evidence.** Routing history is a SHA-256 hash chain a regulator can recompute; the demo tampers with it live and watches the verifier point at the entry.
- **The diode is a first-class actor.** One-way transfer with a counted, visible return-path bounce, quarantine, and a two-person ceremony before a pinned-key verification.
- **One artefact, three stacks, one digest.** Parity turns "consistently deployable everywhere" into a checkable fact.
- **Film and proof from the same engine.** The 3D director and the headless CLI run identical steps.

LIMITATIONS, STATED PLAINLY

- It is a simulator: capacity, cost and latency are modelled, not measured; inference produces a simulated output line.
- The diode is a state machine, not a protocol implementation; real gateways add forward error correction and out-of-band reconciliation.
- Policy attributes are declared by the submitter; in production they come from labelled data sources and an identity provider, not a form.
- Single fictional organisation; coalition releasability across several partners is modelled by one attribute only.
- The architecture is deliberately the production one, so replacing the evaluator with OPA, the signer with cosign and the environments with real clusters changes adapters, not the control room.

WHAT IS NEXT

- **Now:** real adapters behind the same interfaces: OpenAI-compatible endpoints for cloud and on-prem vLLM, cosign-signed OCI bundles, an OPA/Rego export of `rules.json`, Supabase persistence for the ledger.
- **Next:** label-derived attributes from data sources and identity claims; coalition releasability with per-partner enclaves; capacity forecasting and cost-aware scheduling windows; SLSA provenance and SBOMs travelling with the bundle.
- **Later:** an outbound diode for telemetry with differential privacy; hardware attestation of the enclave nodes before load; model watermarking to detect exfiltration; formal checks of the policy (no SECRET path to cloud, ever) with a model checker; a declassification workflow with its own ledger.
- Full roadmap: `docs/ROADMAP.md` in the repository.

CONCLUSION

The brief describes a routing problem, a deployment problem and an evidence problem. MERIDIAN // VANTAGE answers all three with one small engine: policy as data that denies by default, a signed artefact that reaches the enclave only through a one-way link and a two-person check, and a ledger that turns every decision into a verifiable link. It runs from the README, survives a second run and hostile input, and the cinematic control room on top makes the argument in ninety seconds. The next step is not more simulation but real adapters, and the interfaces are already there.